

# ARYAN HUNDIA

Vellore, India | +91 8985470389 | arianhundia125@gmail.com  
linkedin.com/in/aryanhundia | Github

## SUMMARY

AI-security and GRC-focused Computer Science student with hands-on production experience in SOC operations, AI-enabled security workflows, TPRM automation, cloud security, and enterprise compliance at UPL Limited. Built and deployed security infrastructure live in production; multiple national 1st-place wins in cybersecurity competitions judged by CDAC/MeitY, FCRF, and SentinelOne.

## EDUCATION

### Vellore Institute of Technology

B.Tech, Computer and Information Security | CGPA: 9.13

Vellore, Tamil Nadu

2023 – 2027

- **Relevant Coursework:** Information Security, Malware Analysis, Penetration Testing & Vulnerability Analysis, Digital Forensics, Web Application Security, Data Privacy & Protection, Computer Networks, Cloud Architecture & Design, Software Engineering

## EXPERIENCE

### UPL Limited (Fortune 500)

IS Governance Intern – GRC & AI Security

Mumbai, India

Jun 2026 – Jul 2026

- **Security Monitoring & SOC:** Monitored and triaged security events in **Splunk (SIEM)**, analyzing logs for IOCs and anomalous behavior; investigated endpoint threats via **CrowdStrike Falcon (EDR)**; tracked dark web exposure and data leak signals via **CloudSEK** threat intelligence.
- **Risk Assessment:** Conducted an enterprise-wide security risk assessment to identify weak links across UPL's technology infrastructure; mapped identified gaps to control frameworks (ISO 27001, NIST CSF) and Business Continuity Plans; produced risk treatment recommendations adopted by the CISO team.
- **AI-Enabled Security Workflows:** Identified and mitigated **prompt injection** and **data-leakage risks** in production LLM compliance workflows; designed output validation guardrails preventing hallucination from corrupting CISO-level risk scores; aligned controls to ISO 42001 AI governance standards.
- **TPRM & Compliance:** Designed and deployed a live TPRM platform (42 parameters, 40+ vendors) aligned to ISO 27001, ISO 27701, NIST CSF, and SOC 2; authored TPRM policy, conducted end-to-end vendor evaluations (SoA, ITGC/ITAC, SOC 2 review, MSA analysis); ~70% reduction in analyst effort via Python/LLM automation across 12+ compliance processes.
- **Cloud & DLP:** Hardened **Netskope DLP/CASB** policies across cloud access, SaaS, and endpoint channels; validated security requirements within **Microsoft Azure** native tooling (Azure Arc, Power Apps, Power Automate, SharePoint); partnered with cybersecurity and legal stakeholders on IAM and enterprise control implementation.

### Headstarter AI

Software Engineering Fellow

Remote

Jul 2024 – Sep 2024

- Built and shipped five AI-powered projects across a 7-week accelerated fellowship, including an inventory management platform with third-party API integrations, authentication flows; developed under weekly sprint deadlines with production-quality code review standards.

## KEY PROJECTS

### Agent Governance Gateway

Python, FastAPI, MCP, Claude Code Hooks, PowerShell

- Authorization and governance infrastructure for autonomous AI agents: four-tier policy engine (shell, file writes, network, reads) enforcing least-privilege access; human-in-the-loop escalation via PreToolUse hooks – privileged calls intercepted and approved before execution outside agent context.
- Tamper-evident chained audit log, fail-closed error handling, prompt injection containment, privilege escalation prevention, and data exfiltration control; 117 tests passing.

### JARVIS – Voice-Driven Windows Automation Suite

PowerShell, Microsoft Copilot API, Win32 API, Windows SAPI/TTS

- Voice-driven desktop automation using PowerShell and Microsoft Copilot to control native Windows applications (Teams, Outlook) via natural language: composing messages, scheduling meetings, and triaging email threads hands-free.
- Modular command dispatcher mapping Copilot-generated intent to Win32 API calls for system-level control (window management, focus switching, clipboard operations) without GUI interaction; integrated Windows SAPI and custom TTS pipelines for fully hands-free operation.

### TPRM Automation Platform

Next.js 14, Groq LLM, Prisma/SQLite, Recharts, Python, R

- Replaced Excel-based GRC workflow with a production-grade platform: composite risk scoring (Inherent, Residual, Composite), evidence evaluation pipelines, conformity /non-conformity tracking, CISO-level executive dashboards; LLM hallucination guardrails via R statistical validation; live in production at UPL Limited.

**ARGUS – Cyber Forensics Platform** Python, OSINT, Autopsy, AI/ML, FastAPI

- OSINT-integrated digital forensics platform: stylometry analysis, deepfake detection, automated evidence correlation. **1st Place – CDAC/MeitY Cyber Forensics Hackathon (70 teams).**

**PHANTOM TWIN – Behavioral Anomaly Detection** Python, Isolation Forest, React, Groq API

- Profiles users/devices via geo-location, device fingerprints, and peak-hour patterns; routes high-risk sessions into honeypots via Isolation Forest with a sliding-window heuristic engine targeting low-and-slow attacks and credential stuffing.

**LegalShield AI & PII Detection Platform** React, FastAPI, Docker, Kubernetes, OCR, LLMs

- Automated GDPR/DPDP/HIPAA document analysis (99%+ reliability); enterprise PII detection via OCR, NER, and contextual LLM analysis; RBAC, configurable redaction workflows, and async processing pipelines for PII-heavy enterprise environments.

## TECHNICAL SKILLS

|                  |                                                                                                                                                                                                                                     |
|------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| AI Security      | LLM red-teaming, prompt injection detection/containment, AI agent authorization, human-in-the-loop governance, hallucination mitigation, AI output validation, adversarial input testing, AI governance (ISO 42001)                 |
| Cybersecurity    | SOC monitoring, SIEM (Splunk), EDR (CrowdStrike Falcon), CloudSEK TI, IOC/TTP analysis, pen testing, vulnerability assessment, malware analysis, digital forensics, incident response, DLP/CASB (Netskope), Nmap, OWASP ZAP, Nuclei |
| GRC & Compliance | TPRM, risk assessment, IT audit, ITGC/ITAC, ISO 27001/27701/42001, NIST CSF, SOC 2, GDPR, DPDP, HIPAA, BCP risk mapping, MSA review, conformity/non-conformity assessment                                                           |
| Cloud & IAM      | Microsoft Azure (Azure Arc, Power Platform, SharePoint), Netskope CASB, IAM, DLP policy configuration, Docker                                                                                                                       |
| Engineering      | Python, PowerShell, JavaScript, FastAPI, Next.js, React.js, PostgreSQL, MongoDB, Git, MCP, REST APIs                                                                                                                                |

## CERTIFICATIONS

- Microsoft Certified: Azure AI Apps and Agents Developer Associate
- Microsoft Certified: Cloud and AI Security Engineer Associate
- OCI 2025 Generative AI Professional – Oracle (*ongoing*)
- OCI 2025 Data Science Professional – Oracle
- Machine Learning – Columbia University
- Industrial Cybersecurity Essentials – Cisco | Introduction to Cybersecurity – Cisco
- GRC Fundamentals – CyberExam

## PATENTS FILED

- **Hardware Data-Diode Apparatus for Prompt-Injection-Resilient LLM Malware Analysis** – physical relay gating with canary-token verification; 96.6% outbound-action containment rate vs. 59.0% software-only baseline
- **Hardware Sampling-Firewall for Steganographic Covert-Channel Denial in LLM Output** – TRNG-keyed perturbation on inference accelerator output bus; 82.1% mean covert-channel bandwidth reduction, scheme-agnostic
- **HSM-Gated Action-Graph Conformance for Browser-Agent Session Hijack Interdiction** – FPGA-resident scoring with physical network relay; 95.8% hijack detection accuracy, 5.35 ms gating latency (8.8x faster than software)
- **Hardware-Traced Circuit Breaker for AI-Hallucinated Dependency Containment** – instruction-pointer tagging of slopsquatted packages with physical uplink disconnect; 96.2% detection accuracy, 17.1x latency advantage
- **Hardware-Timestamped Challenge-Response for AI Agent Integrity Verification** – PTP-disciplined behavioral confusion scoring with FPGA relay quarantine; physical disconnection in 4.33 ms of compromise detection

## HACKATHONS & RECOGNITION

- **1st Place** – FCRF National Cybercrime Hackathon, FutureCrime Summit 2026, Bharat Mandapam, New Delhi
- **1st Place** – CDAC/MeitY Cyber Forensics Hackathon (70 teams)
- **1st Place** – AI Hackathon Challenge, VIT
- **3rd Place** – ThreatOps National, FINSEC’26 (SentinelOne)
- **Top 3** – PULSE of Viksit Maharashtra 2026